

# Customer churn

## Business Objective:

Customer churn is a big problem for telecommunications companies. Indeed, their annual churn rates are usually higher than 10%. For that reason, they develop strategies to keep as many clients as possible. This is a classification project since the variable to be predicted is binary (churn or loyal customer). The goal here is to model churn probability, conditioned on the customer features.

## Data Set Details:

Each row corresponds to a client of a telecommunications company for whom it has collected information about the type of plan they have contracted, the minutes they have talked, or the charge they pay every month.

The data set includes the following variables:

- state: Categorical, for the 51 states and the District of Columbia.
- Area.code
- account.length: how long the account has been active.
- voice.plan: yes or no, voicemail plan.
- voice.messages: number of voicemail messages.
- intl.plan: yes or no, international plan.
- intl.mins: minutes customer used service to make international calls.
- intl.calls: total number of international calls.
- intl.charge: total international charge.
- day.mins: minutes customer used service during the day.
- day.calls: total number of calls during the day.
- day.charge: total charge during the day.
- eve.mins: minutes customer used service during the evening.
- eve.calls: total number of calls during the evening.
- eve.charge: total charge during the evening.
- night.mins: minutes customer used service during the night.
- night.calls: total number of calls during the night.
- night.charge: total charge during the night.
- customer.calls: number of calls to customer service.
- churn: Categorical, yes or no. Indicator of whether the customer has left the company (yes or no).

In [1]:

```
#Importing Neccesary Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings

%matplotlib inline
warnings.filterwarnings("ignore")
```

In [2]:

```
#Importing dataset
data = pd.read_csv('Churn.csv')
```

In [3]:

```
data.head()
```

Out[3]:

	Unnamed: 0	state	area.code	account.length	voice.plan	voice.messages	intl.plan	int
0	1	KS	area_code_415	128	yes	25	no	
1	2	OH	area_code_415	107	yes	26	no	
2	3	NJ	area_code_415	137	no	0	no	
3	4	OH	area_code_408	84	no	0	yes	
4	5	OK	area_code_415	75	no	0	yes	

5 rows × 21 columns

In [4]:

```
df=data.drop(['Unnamed: 0'], axis=1)
```

In [5]:

```
df.head()
```

Out[5]:

	state	area.code	account.length	voice.plan	voice.messages	intl.plan	intl.mins	intl.c
0	KS	area_code_415	128	yes	25	no	10.0	
1	OH	area_code_415	107	yes	26	no	13.7	
2	NJ	area_code_415	137	no	0	no	12.2	
3	OH	area_code_408	84	no	0	yes	6.6	
4	OK	area_code_415	75	no	0	yes	10.1	

In [6]:

```
#Rows and columns  
df.shape
```

Out[6]:

(5000, 20)

In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5000 entries, 0 to 4999  
Data columns (total 20 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   state                 5000 non-null   object  
1   area.code             5000 non-null   object  
2   account.length       5000 non-null   int64  
3   voice.plan            5000 non-null   object  
4   voice.messages       5000 non-null   int64  
5   intl.plan            5000 non-null   object  
6   intl.mins             5000 non-null   float64  
7   intl.calls           5000 non-null   int64  
8   intl.charge          5000 non-null   float64  
9   day.mins             5000 non-null   float64  
10  day.calls            5000 non-null   int64  
11  day.charge           5000 non-null   object  
12  eve.mins             5000 non-null   object  
13  eve.calls            5000 non-null   int64  
14  eve.charge           5000 non-null   float64  
15  night.mins           5000 non-null   float64  
16  night.calls          5000 non-null   int64  
17  night.charge         5000 non-null   float64  
18  customer.calls       5000 non-null   int64  
19  churn                5000 non-null   object  
dtypes: float64(6), int64(7), object(7)  
memory usage: 781.4+ KB
```

In [8]:

```
df.dtypes
```

Out[8]:

```
state          object
area.code      object
account.length int64
voice.plan     object
voice.messages int64
intl.plan     object
intl.mins     float64
intl.calls    int64
intl.charge   float64
day.mins      float64
day.calls     int64
day.charge    object
eve.mins      object
eve.calls     int64
eve.charge    float64
night.mins    float64
night.calls   int64
night.charge  float64
customer.calls int64
churn         object
dtype: object
```

## Report

- we have 2 features with wrong data type which is day.charge and eve.mins

In [9]:

```
df['day.charge'] = df['day.charge'].astype(float)
```

In [10]:

```
df['eve.mins'] = df['eve.mins'].astype(float)
```

In [11]:

```
df.dtypes
```

Out[11]:

```
state          object
area.code      object
account.length int64
voice.plan     object
voice.messages int64
intl.plan     object
intl.mins     float64
intl.calls    int64
intl.charge   float64
day.mins      float64
day.calls     int64
day.charge    float64
eve.mins      float64
eve.calls     int64
eve.charge    float64
night.mins    float64
night.calls   int64
night.charge  float64
customer.calls int64
churn         object
dtype: object
```

In [12]:

```
#Extracting cat feature from the dataset
cat_feature = [feature for feature in df.columns if df[feature].dtypes == 'O']
```

In [13]:

```
cat_feature
```

Out[13]:

```
['state', 'area.code', 'voice.plan', 'intl.plan', 'churn']
```

In [14]:

```
#Extracting cat feature from the dataset
num_feature = [feature for feature in df.columns if df[feature].dtypes != 'O']
```

In [15]:

```
num_feature
```

Out[15]:

```
['account.length',  
 'voice.messages',  
 'intl.mins',  
 'intl.calls',  
 'intl.charge',  
 'day.mins',  
 'day.calls',  
 'day.charge',  
 'eve.mins',  
 'eve.calls',  
 'eve.charge',  
 'night.mins',  
 'night.calls',  
 'night.charge',  
 'customer.calls']
```

In [16]:

```
#checking null values if there are any  
df.isnull().sum()
```

Out[16]:

```
state                0  
area.code            0  
account.length      0  
voice.plan          0  
voice.messages      0  
intl.plan           0  
intl.mins           0  
intl.calls          0  
intl.charge         0  
day.mins            0  
day.calls           0  
day.charge          7  
eve.mins            24  
eve.calls           0  
eve.charge          0  
night.mins          0  
night.calls         0  
night.charge        0  
customer.calls      0  
churn               0  
dtype: int64
```

In [17]:

```
#Extracting unique value from the dataset
```

```
for feature in df.columns:
```

```
    print(f"feature {feature} has {df[feature].unique()} unique values \n")
```

```
feature state has ['KS' 'OH' 'NJ' 'OK' 'AL' 'MA' 'MO' 'LA' 'WV' 'IN' 'RI'
'IA' 'MT' 'NY'
'ID' 'VT' 'VA' 'TX' 'FL' 'CO' 'AZ' 'SC' 'NE' 'WY' 'HI' 'IL' 'NH' 'GA'
'AK' 'MD' 'AR' 'WI' 'OR' 'MI' 'DE' 'UT' 'CA' 'MN' 'SD' 'NC' 'WA' 'NM'
'NV' 'DC' 'KY' 'ME' 'MS' 'TN' 'PA' 'CT' 'ND'] unique values
```

```
feature area.code has ['area_code_415' 'area_code_408' 'area_code_510'] unique values
```

```
feature account.length has [128 107 137 84 75 118 121 147 117 141 65 7
4 168 95 62 161 85 93
76 73 77 130 111 132 174 57 54 20 49 142 172 12 72 36 78 136
149 98 135 34 160 64 59 119 97 52 60 10 96 87 81 68 125 116
38 40 43 113 126 150 138 162 90 50 82 144 46 70 55 106 94 155
80 104 99 120 108 122 157 103 63 112 41 193 61 92 131 163 91 127
110 140 83 145 56 151 139 6 115 146 185 148 32 25 179 67 19 170
164 51 208 53 105 66 86 35 88 123 45 100 215 22 33 114 24 101
143 48 71 167 89 199 166 158 196 209 16 39 173 129 44 79 31 124
37 159 194 154 21 133 224 58 11 109 102 165 18 30 176 47 190 152
26 69 186 171 28 153 169 13 27 3 42 189 156 134 243 23 1 205
200 5 9 178 181 182 217 177 210 29 180 2 17 7 212 232 192 195
197 225 184 191 201 15 183 202 8 175 4 188 204 221 187 14 238 216
222 233] unique values
```

```
feature voice.plan has ['yes' 'no'] unique values
```

```
feature voice.messages has [25 26 0 24 37 27 33 39 30 41 28 34 46 29 35 2
1 32 42 36 22 23 43 31 38
40 48 18 17 45 16 20 14 19 51 15 11 12 47 8 44 49 4 10 13 50 9 6 52]
unique values
```

```
feature intl.plan has ['no' 'yes'] unique values
```

```
feature intl.mins has [10. 13.7 12.2 6.6 10.1 6.3 7.5 7.1 8.7 11.2 1
2.7 9.1 12.3 13.1
5.4 13.8 8.1 13. 10.6 5.7 9.5 7.7 10.3 15.5 14.7 11.1 14.2 12.6
11.8 8.3 14.5 10.5 9.4 14.6 9.2 3.5 8.5 13.2 7.4 8.8 11. 7.8
6.8 11.4 9.3 9.7 10.2 8. 5.8 12.1 12. 11.6 8.2 6.2 7.3 6.1
11.7 15. 9.8 12.4 8.6 10.9 13.9 8.9 7.9 5.3 4.4 12.5 11.3 9.
9.6 13.3 20. 7.2 6.4 14.1 14.3 6.9 11.5 15.8 12.8 16.2 0. 11.9
9.9 8.4 10.8 13.4 10.7 17.6 4.7 2.7 13.5 12.9 14.4 10.4 6.7 15.4
4.5 6.5 15.6 5.9 18.9 7.6 5. 7. 14. 18. 16. 14.8 3.7 2.
4.8 15.3 6. 13.6 17.2 17.5 5.6 18.2 3.6 16.5 4.6 5.1 4.1 16.3
14.9 16.4 16.7 1.3 15.2 15.1 15.9 5.5 16.1 4. 16.9 5.2 4.2 15.7
17. 3.9 3.8 2.2 17.1 4.9 17.9 17.3 18.4 17.8 4.3 2.9 3.1 3.3
2.6 3.4 1.1 18.3 16.6 2.1 2.4 2.5 18.7 16.8 0.4 19.3 19.2 19.7
18.5 17.7] unique values
```

```
feature intl.calls has [3 5 7 6 4 2 9 19 1 10 15 8 11 0 12 13 18
14 16 20 17] unique values
```

```
feature intl.charge has [2.7 3.7 3.29 1.78 2.73 1.7 2.03 1.92 2.35 3.02
3.43 2.46 3.32 3.54
1.46 3.73 2.19 3.51 2.86 1.54 2.57 2.08 2.78 4.19 3.97 3. 3.83 3.4
3.19 2.24 3.92 2.84 2.54 3.94 2.48 0.95 2.3 3.56 2. 2.38 2.97 2.11
1.84 3.08 2.51 2.62 2.75 2.16 1.57 3.27 3.24 3.13 2.21 1.67 1.97 1.65
3.16 4.05 2.65 3.35 2.32 2.94 3.75 2.4 2.13 1.43 1.19 3.38 3.05 2.43
2.59 3.59 5.4 1.94 1.73 3.81 3.86 1.86 3.11 4.27 3.46 4.37 0. 3.21
2.67 2.27 2.92 3.62 2.89 4.75 1.27 0.73 3.65 3.48 3.89 2.81 1.81 4.16
1.22 1.76 4.21 1.59 5.1 2.05 1.35 1.89 3.78 4.86 4.32 4. 1. 0.54
1.3 4.13 1.62 3.67 4.64 4.73 1.51 4.91 0.97 4.46 1.24 1.38 1.11 4.4
```



```

4.02 4.43 4.51 0.35 4.1 4.08 4.29 1.49 4.35 1.08 4.56 1.4 1.13 4.24
4.59 1.05 1.03 0.59 4.62 1.32 4.83 4.67 4.97 4.81 1.16 0.78 0.84 0.89
0.7 0.92 0.3 4.94 4.48 0.57 0.65 0.68 5.05 4.54 0.11 5.21 5.18 5.32
5. 4.78] unique values

```

```

feature day.mins has [265.1 161.6 243.4 ... 188.7 7.2 170. ] unique values

```

```

feature day.calls has [110 123 114 71 113 98 88 79 97 84 137 127 96
70 67 139 66 90
117 89 112 103 86 76 115 73 109 95 105 121 118 94 80 128 64 106
102 85 82 77 120 133 135 108 57 83 129 91 92 74 93 101 146 72
99 104 125 61 100 87 131 65 124 119 52 68 107 47 116 151 126 122
111 145 78 136 140 148 81 55 69 158 134 130 63 53 75 141 163 59
132 138 54 58 62 144 143 147 36 40 150 56 51 165 30 48 60 42
0 45 160 149 152 142 156 35 49 157 44 50 34 39 46] unique values

```

```

feature day.charge has [45.07 27.47 41.38 ... 32.08 1.22 28.9 ] unique values

```

```

feature eve.mins has [197.4 195.5 121.2 ... 302.3 280.6 340.3] unique values

```

```

feature eve.calls has [ 99 103 110 88 122 101 108 94 80 111 83 148 71
75 76 97 90 65
93 121 102 72 112 100 84 109 63 107 115 119 116 92 85 98 118 74
117 58 96 66 67 62 77 164 126 142 64 104 79 95 86 105 81 113
106 59 48 82 87 123 114 140 128 60 78 125 91 46 138 129 89 133
136 57 135 139 51 70 151 137 134 73 152 168 68 120 69 127 132 143
61 124 42 54 131 52 149 56 37 130 49 146 147 55 12 50 157 155
45 144 36 156 53 141 44 153 154 150 43 0 145 159 170 47 169 38]
unique values

```

```

feature eve.charge has [16.78 16.62 10.3 ... 25.7 23.85 28.93] unique values

```

```

feature night.mins has [244.7 254.4 162.6 ... 75.1 297.5 224.4] unique values

```

```

feature night.calls has [ 91 103 104 89 121 118 96 90 97 111 94 128 1
15 99 75 108 74 133
64 78 105 68 102 148 98 116 71 109 107 135 92 86 127 79 87 129
57 77 95 54 106 53 67 139 60 100 61 73 113 76 119 88 84 62
137 72 142 114 126 122 81 123 117 82 80 120 130 134 59 112 132 110
101 150 69 131 83 93 124 136 125 66 143 58 55 85 56 70 46 42
152 44 145 50 153 49 175 63 138 154 140 141 146 65 51 151 158 155
157 147 144 149 166 52 33 156 38 36 48 164 40 168 161 159 160 170
41 12 165 43 0] unique values

```

```

feature night.charge has [11.01 11.45 7.32 ... 4.65 3.65 3.38] unique values

```

```

feature customer.calls has [1 0 2 3 4 5 7 9 6 8] unique values

```

```

feature churn has ['no' 'yes'] unique values

```

In [18]:

```
df.columns
```

Out[18]:

```
Index(['state', 'area.code', 'account.length', 'voice.plan', 'voice.messages',  
      'intl.plan', 'intl.mins', 'intl.calls', 'intl.charge', 'day.mins',  
      'day.calls', 'day.charge', 'eve.mins', 'eve.calls', 'eve.charge',  
      'night.mins', 'night.calls', 'night.charge', 'customer.calls', 'churn'],  
      dtype='object')
```

In [19]:

```
df[df['state']=='IA'].head(50)
```

Out[19]:

	state	area.code	account.length	voice.plan	voice.messages	intl.plan	intl.mins	ii
12	IA	area_code_408	168	no	0	no	11.2	
14	IA	area_code_415	62	no	0	no	13.1	
50	IA	area_code_408	52	no	0	no	7.8	
100	IA	area_code_510	98	yes	21	no	4.4	
150	IA	area_code_408	113	no	0	no	10.6	
162	IA	area_code_510	141	yes	36	no	10.5	
227	IA	area_code_408	126	yes	27	no	9.6	
303	IA	area_code_415	158	no	0	no	9.1	
328	IA	area_code_510	76	no	0	no	12.5	
489	IA	area_code_415	130	no	0	no	12.3	
699	IA	area_code_415	98	no	0	no	11.0	
855	IA	area_code_510	66	no	0	no	6.2	
887	IA	area_code_408	128	no	0	no	7.6	
912	IA	area_code_510	45	no	0	no	10.6	
941	IA	area_code_510	63	no	0	no	6.5	
1113	IA	area_code_415	152	no	0	no	10.6	
1175	IA	area_code_415	134	yes	32	no	8.6	
1206	IA	area_code_510	92	yes	25	no	9.7	
1234	IA	area_code_408	86	no	0	no	9.8	
1266	IA	area_code_415	42	no	0	no	9.0	
1300	IA	area_code_510	46	no	0	no	10.2	
1356	IA	area_code_415	118	no	0	no	11.3	
1494	IA	area_code_415	129	no	0	no	4.9	
1527	IA	area_code_510	36	no	0	no	9.0	
1530	IA	area_code_510	81	no	0	no	8.0	
1605	IA	area_code_415	73	no	0	no	11.1	
1838	IA	area_code_408	1	yes	26	no	8.1	
2031	IA	area_code_510	130	no	0	no	11.4	
2035	IA	area_code_510	81	no	0	no	8.7	
2040	IA	area_code_510	105	yes	15	no	9.7	
2085	IA	area_code_415	75	no	0	no	11.1	
2149	IA	area_code_415	120	yes	33	no	11.6	
2261	IA	area_code_408	100	no	0	no	9.4	
2416	IA	area_code_510	113	no	0	no	11.9	
2561	IA	area_code_510	143	yes	33	no	7.7	
2570	IA	area_code_415	64	yes	43	no	8.5	
2632	IA	area_code_415	89	yes	35	no	11.8	

	state	area.code	account.length	voice.plan	voice.messages	intl.plan	intl.mins	ii
2661	IA	area_code_415	197	no	0	no	9.5	
2677	IA	area_code_415	44	no	0	no	7.3	
2783	IA	area_code_415	79	yes	17	no	9.1	
2856	IA	area_code_415	123	no	0	no	10.0	
3125	IA	area_code_510	40	no	0	no	10.5	
3193	IA	area_code_415	88	no	0	no	10.8	
3308	IA	area_code_415	45	no	0	no	13.3	
3338	IA	area_code_415	117	no	0	no	6.9	
3365	IA	area_code_415	92	no	0	no	10.0	
3424	IA	area_code_408	82	no	0	no	10.3	
3432	IA	area_code_408	74	no	0	no	9.0	
3476	IA	area_code_408	93	no	0	no	12.6	
3481	IA	area_code_510	157	no	0	no	13.6	

In [20]:

```
# Observation in few columns we have anonymous value "Nan" we will replace it with np.nan
```

In [21]:

```
#Basic stats about the dataset
df.describe().T
```

Out[21]:

	count	mean	std	min	25%	50%	75%	max
account.length	5000.0	100.258600	39.694560	1.0	73.000	100.00	127.00	243.00
voice.messages	5000.0	7.755200	13.546393	0.0	0.000	0.00	17.00	52.00
intl.mins	5000.0	10.261780	2.761396	0.0	8.500	10.30	12.00	20.00
intl.calls	5000.0	4.435200	2.456788	0.0	3.000	4.00	6.00	20.00
intl.charge	5000.0	2.771196	0.745514	0.0	2.300	2.78	3.24	5.40
day.mins	5000.0	180.288900	53.894699	0.0	143.700	180.10	216.20	351.50
day.calls	5000.0	100.029400	19.831197	0.0	87.000	100.00	113.00	165.00
day.charge	4993.0	30.653501	9.166356	0.0	24.430	30.62	36.75	59.76
eve.mins	4976.0	200.580326	50.554637	0.0	166.275	201.00	234.10	363.70
eve.calls	5000.0	100.191000	19.826496	0.0	87.000	100.00	114.00	170.00
eve.charge	5000.0	17.054322	4.296843	0.0	14.140	17.09	19.90	30.91
night.mins	5000.0	200.391620	50.527789	0.0	166.900	200.40	234.70	395.00
night.calls	5000.0	99.919200	19.958686	0.0	87.000	100.00	113.00	175.00
night.charge	5000.0	9.017732	2.273763	0.0	7.510	9.02	10.56	17.77
customer.calls	5000.0	1.570400	1.306363	0.0	1.000	1.00	2.00	9.00

In [22]:

```
#Checking unique values for target variable  
df['churn'].value_counts()
```

Out[22]:

```
no      4293  
yes      707  
Name: churn, dtype: int64
```

In [23]:

```
100*df['churn'].value_counts()/len(data['churn'])
```

Out[23]:


```
no      85.86  
yes     14.14  
Name: churn, dtype: float64
```

### Visualization of target feature on unique values

In [24]:

```
yes = df[df['churn']=='yes'].shape[0]
no = df[df['churn']=='no'].shape[0]
print("No: " + str(no) + ", yes: " + str(yes))
sns.catplot(data=df, x="churn", kind="count", palette="winter_r", alpha=.6)
plt.show()
```

yes: 4293, yes: 707



## Report

The target feature are highly imbalanced

Class imbalance is a scenario that arises when we have unequal distribution of class in a dataset i.e. the no. of data points in the no churn (majority class) very large compared to that of the yes churn (minority class)

If the imbalanced data is not treated beforehand, then this will degrade the performance of the classifier model.

Hence we should handle imbalanced data with certain methods.

## How to handle Imbalance Data ?

Resampling data is one of the most commonly preferred approaches to deal with an imbalanced dataset.

There are broadly two types of methods for this i) Undersampling ii) Oversampling. In most cases, oversampling is preferred over undersampling techniques. The reason being, in undersampling we tend to remove instances from data that may be carrying some important information.

## SMOTE: Synthetic Minority Oversampling Technique

SMOTE is an oversampling technique where the synthetic samples are generated for the minority class.

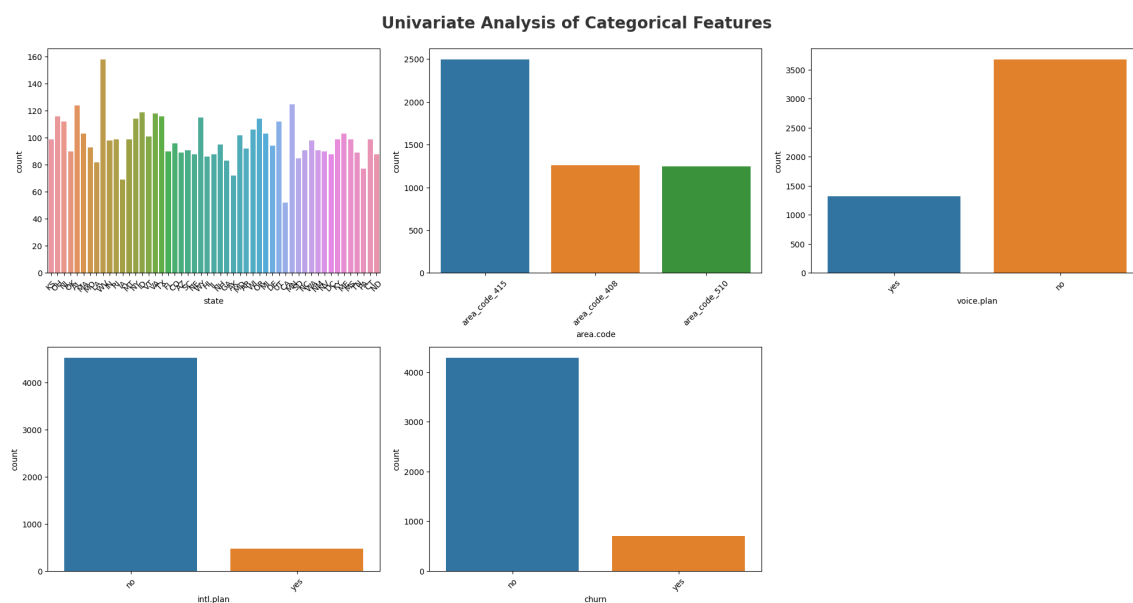
Hybridization techniques involve combining both undersampling and oversampling techniques. This is done to optimize the performance of classifier models for the samples created as part of these techniques.

It only duplicates the data and it won't add and new information. Hence we look at some different techniques.

## Visulization of Categorical feature

In [25]:

```
# categorical columns
plt.figure(figsize=(20, 15))
plt.suptitle('Univariate Analysis of Categorical Features', fontsize=20, fontweight='bold')
for i in range(0, len(cat_feature)):
    plt.subplot(3, 3, i+1)
    sns.countplot(x=df[cat_feature[i]])
    plt.xlabel(cat_feature[i])
    plt.xticks(rotation=45)
    plt.tight_layout()
```





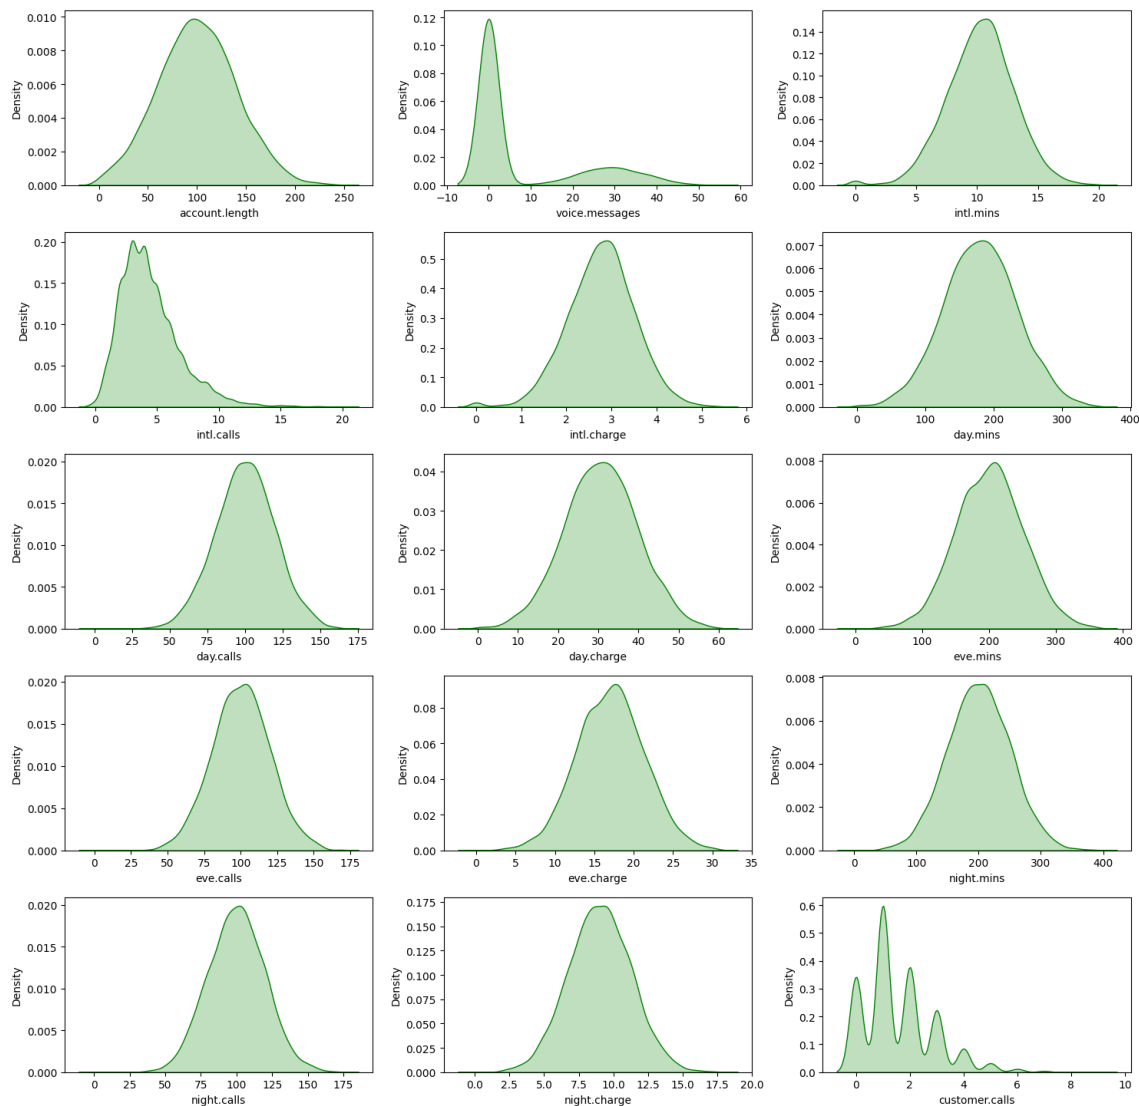
# Visualization of numerical features

In [26]:

```
plt.figure(figsize=(15, 15))
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20, fontweight='bold')

for i in range(0, len(num_feature)):
    plt.subplot(5, 3, i+1)
    sns.kdeplot(x=df[num_feature[i]], shade=True, color='g')
    plt.xlabel(num_feature[i])
    plt.tight_layout()
```

Univariate Analysis of Numerical Features



## Report

Most of the features are normally distributed only few features are little bit skewed towards right or left

## Check Multicollinearity in Numerical features

In [27]:


```
df.corr()
```

Out[27]:

	account.length	voice.messages	intl.mins	intl.calls	intl.charge	day.mins
account.length	1.000000	-0.014575	0.001291	0.014277	0.001292	-0.001017
voice.messages	-0.014575	1.000000	0.002463	0.000124	0.002505	0.005381
intl.mins	0.001291	0.002463	1.000000	0.016791	0.999993	-0.019486
intl.calls	0.014277	0.000124	0.016791	1.000000	0.016900	-0.001303
intl.charge	0.001292	0.002505	0.999993	0.016900	1.000000	-0.019415
day.mins	-0.001017	0.005381	-0.019486	-0.001303	-0.019415	1.000000
day.calls	0.028240	0.000883	0.013097	0.010893	0.013161	0.001935
day.charge	-0.000800	0.005140	-0.019295	-0.001599	-0.019225	1.000000
eve.mins	-0.008706	0.018917	-0.000347	0.008001	-0.000324	-0.010931
eve.calls	0.009143	-0.003954	-0.007458	0.005574	-0.007507	0.008128
eve.charge	-0.009587	0.019496	0.000132	0.008393	0.000155	-0.010760
night.mins	0.000668	0.005541	-0.006721	-0.017214	-0.006655	0.011799
night.calls	-0.007825	0.002676	0.000391	-0.000156	0.000368	0.004236
night.charge	0.000656	0.005535	-0.006717	-0.017182	-0.006650	0.011783
customer.calls	-0.001445	-0.007086	-0.012122	-0.019147	-0.012180	0.002733

In [28]:

```
plt.figure(figsize = (15,10))
sns.heatmap(df.corr(), cmap="CMRmap", annot=True)
plt.show()
```



## Report

feature customer.calls is negatively correlated with feature voice.messages, intl.min, intl.call, day.call, nights.min, night.charge, night.calls

feature intl.charge and feature intl.min, feature night.mins and night.charge are highly correlated likewise most of the features are negative correlated

In [29]:

```
Top_10_state = df['state'].value_counts().head(10)
Top_10_state
```

Out[29]:

```
WV    158
MN    125
AL    124
ID    119
VA    118
OH    116
TX    116
WY    115
NY    114
OR    114
Name: state, dtype: int64
```

In [30]:


```
Top_10_state_intlcalls = Top_10_state.index
Top_10_state_intlcalls
```

Out[30]:

```
Index(['WV', 'MN', 'AL', 'ID', 'VA', 'OH', 'TX', 'WY', 'NY', 'OR'], dtype
='object')
```

In [31]:

```
plt.subplots(figsize=(14,7))
sns.countplot(x="state", data=df, order = Top_10_state_intlcalls)
plt.title("Top 10 state by most international calls", weight="bold", fontsize=20, pad=20)
plt.ylabel("Count", weight="bold", fontsize=20)
plt.xlabel("state", weight="bold", fontsize=16)
plt.xticks(rotation= 45)
plt.xlim(-1,10.5)
plt.show()
```



**Check the mean international calls for state WV which have highest number of international calls**

In [32]:

```
WV = df[df['state'] == 'WV']['intl.calls'].mean()
print(f"state WV has average {round(WV)} international calls every month ")
```

```
state WV has average 4 international calls every month
```

## Report

- As per the graph state WV has the most no of international calls among all other states
- Following WV we have MN and AL
- Mean International calls from State WV are 4 per month

In [33]:

```
df.columns
```

Out[33]:

```
Index(['state', 'area.code', 'account.length', 'voice.plan', 'voice.messages',  
      'intl.plan', 'intl.mins', 'intl.calls', 'intl.charge', 'day.mins',  
      'day.calls', 'day.charge', 'eve.mins', 'eve.calls', 'eve.charge',  
      'night.mins', 'night.calls', 'night.charge', 'customer.calls', 'churn'],  
      dtype='object')
```

In [34]:

```
International_details = df.groupby('state')['intl.mins', 'intl.calls', 'intl.charge'].su  
International_details
```

Out[34]:

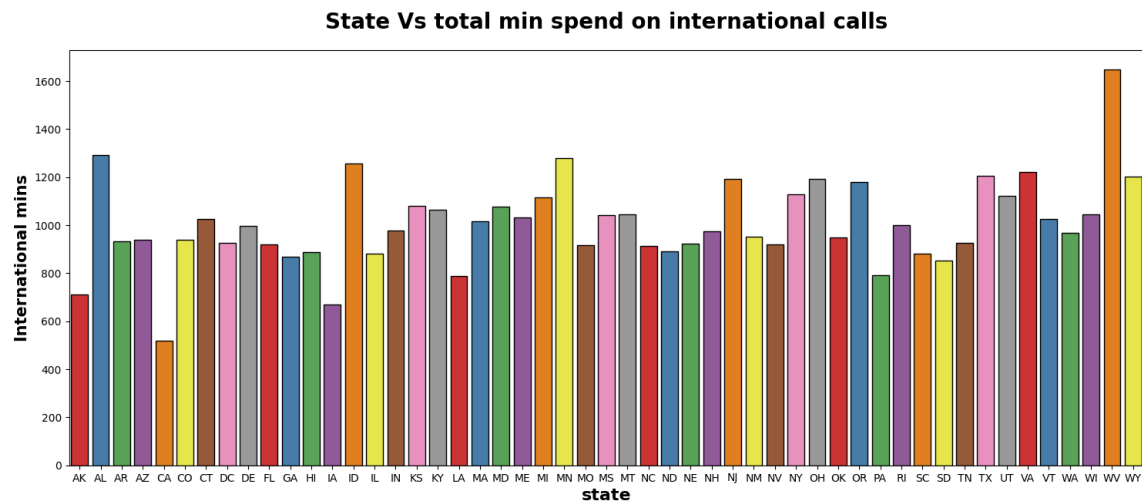
	state	intl.mins	intl.calls	intl.charge
0	AK	711.5	335	192.15
1	AL	1291.2	558	348.65
2	AR	933.9	437	252.16
3	AZ	938.5	423	253.44
4	CA	518.7	226	140.09
5	CO	938.8	416	253.52
6	CT	1027.2	403	277.35
7	DC	925.3	349	249.90
8	DE	997.0	397	269.31
9	FL	920.0	382	248.44
10	GA	869.3	358	234.77
11	HI	889.0	420	240.08
12	IA	670.6	291	181.10
13	ID	1256.6	553	339.35
14	IL	883.0	347	238.46
15	IN	979.0	400	264.36
16	KS	1080.8	437	291.84
17	KY	1063.0	405	287.05
18	LA	789.2	378	213.18
19	MA	1016.5	475	274.53
20	MD	1076.3	452	290.65
21	ME	1032.7	449	278.90
22	MI	1115.7	491	301.25
23	MN	1279.1	526	345.42
24	MO	917.4	465	247.75
25	MS	1041.9	429	281.38
26	MT	1046.1	452	282.48
27	NC	912.5	400	246.42
28	ND	892.8	415	241.11
29	NE	922.9	361	249.25
30	NH	973.5	423	262.91
31	NJ	1192.3	511	322.01
32	NM	951.7	422	257.07
33	NV	920.8	391	248.61
34	NY	1127.2	519	304.34
35	OH	1191.5	494	321.76
36	OK	947.7	419	255.99

	state	intl.mins	intl.calls	intl.charge
37	OR	1180.6	483	318.83
38	PA	791.5	309	213.69
39	RI	999.9	433	270.01
40	SC	881.9	381	238.18
41	SD	851.9	393	230.05
42	TN	927.2	395	250.38
43	TX	1205.0	492	325.38
44	UT	1121.1	515	302.72
45	VA	1220.4	564	329.57
46	VT	1024.6	486	276.69
47	WA	969.0	450	261.71
48	WI	1045.7	431	282.38
49	WV	1647.3	701	444.88
50	WY	1201.6	534	324.48

In [35]:

```
plt.subplots(figsize=(18,7))
sns.barplot(x=International_details.state, y=International_details['intl.mins'],ec = "b")
plt.title("State Vs total min spend on international calls", weight="bold",fontsize=20,
plt.ylabel("International mins", weight="bold", fontsize=15)
plt.xlabel("state", weight="bold", fontsize=16)

plt.show()
```



## Report

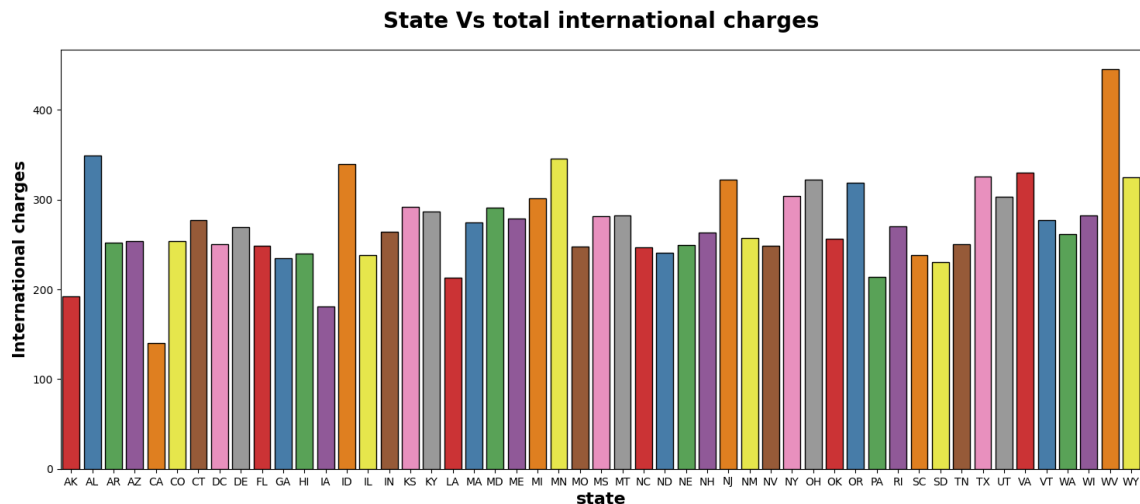
- state WV has total above 1600 mins spend on international calls all over time



In [36]:

```
plt.subplots(figsize=(18,7))
sns.barplot(x=International_details.state, y=International_details['intl.charge'],ec = "
plt.title("State Vs total international charges", weight="bold",fontsize=20, pad=20)
plt.ylabel("International charges", weight="bold", fontsize=15)
plt.xlabel("state", weight="bold", fontsize=16)

plt.show()
```



## Report


- Its obvious that those state which spent more time on international calls have more charge rate which is state 'WV'

In [37]:

```
plt.subplots(figsize=(18,7))
sns.scatterplot(x=df['area.code'], y=df['intl.mins'], hue=df['voice.plan'])
```

Out[37]:

<AxesSubplot: xlabel='area.code', ylabel='intl.mins'>



Type *Markdown* and LaTeX:  $\alpha^2$

## Report

- Higher the total duration of call during international call greater the international charge

In [39]:

```
df['TotalDaysMins'] = df['day.mins'] + df['eve.mins'] + df['night.mins']
```

In [40]:


```
df['TotalCharge'] = df['day.charge'] + df['eve.charge'] + df['night.charge']
```

In [41]:

```
df['TotalDaysCalls'] = df['day.calls'] + df['eve.calls'] + df['night.calls']
```

In [45]:

```
plt.subplots(figsize=(20,7))
for feature in num_feature:
    df.boxplot(column=num_feature, grid=False)
```



## Report


- The no. of 'no' from class churn is greater than no of 'yes', which means loyal customers are higher than churn
- Most of the loyal customer purchased the voice plan
- Rate of Loyal customers is higher than churn customers

In [47]:

```
plt.subplots(figsize=(18,7))  
sns.histplot(data=df, x="state", hue="voice.plan", multiple="stack")
```

Out[47]:

<AxesSubplot: xlabel='state', ylabel='Count'>



## Report

- We can say no of customers with voice plan is less compare to customers who have voice plan

## Feature Engineering

In [48]:

```
cat_feature
```

Out[48]:

```
['state', 'area.code', 'voice.plan', 'intl.plan', 'churn']
```

In [49]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 23 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   state                 5000 non-null   object
 1   area.code             5000 non-null   object
 2   account.length       5000 non-null   int64
 3   voice.plan           5000 non-null   object
 4   voice.messages       5000 non-null   int64
 5   intl.plan            5000 non-null   object
 6   intl.mins            5000 non-null   float64
 7   intl.calls           5000 non-null   int64
 8   intl.charge          5000 non-null   float64
 9   day.mins             5000 non-null   float64
10   day.calls            5000 non-null   int64
11   day.charge           4993 non-null   float64
12   eve.mins             4976 non-null   float64
13   eve.calls            5000 non-null   int64
14   eve.charge           5000 non-null   float64
15   night.mins           5000 non-null   float64
16   night.calls          5000 non-null   int64
17   night.charge         5000 non-null   float64
18   customer.calls       5000 non-null   int64
19   churn                5000 non-null   object
20   TotalDaysMins        4976 non-null   float64
21   TotalCharge          4993 non-null   float64
22   TotalDaysCalls       5000 non-null   int64
dtypes: float64(10), int64(8), object(5)
memory usage: 898.6+ KB
```

In [50]:

```
df.isnull().sum()
```

Out[50]:

```
state                0
area.code            0
account.length      0
voice.plan          0
voice.messages      0
intl.plan           0
intl.mins           0
intl.calls          0
intl.charge         0
day.mins            0
day.calls           0
day.charge          7
eve.mins           24
eve.calls           0
eve.charge          0
night.mins          0
night.calls         0
night.charge        0
customer.calls      0
churn               0
TotalDaysMins       24
TotalCharge          7
TotalDaysCalls      0
dtype: int64
```

## Filling null values

In [51]:

```
df['TotalDaysMins'] = df['TotalDaysMins'].fillna(df['TotalDaysMins'].mean())
```

In [52]:

```
df['TotalCharge'] = df['TotalCharge'].fillna(df['TotalCharge'].mean())
```

In [53]:

```
df.isnull().sum()
```

Out[53]:

```
state          0
area.code      0
account.length 0
voice.plan     0
voice.messages 0
intl.plan      0
intl.mins      0
intl.calls     0
intl.charge    0
day.mins       0
day.calls      0
day.charge     7
eve.mins       24
eve.calls      0
eve.charge     0
night.mins     0
night.calls    0
night.charge   0
customer.calls 0
churn          0
TotalDaysMins  0
TotalCharge    0
TotalDaysCalls 0
dtype: int64
```

In [54]:

```
df
```

Out[54]:

	state	area.code	account.length	voice.plan	voice.messages	intl.plan	intl.mins	ii
0	KS	area_code_415	128	yes	25	no	10.0	
1	OH	area_code_415	107	yes	26	no	13.7	
2	NJ	area_code_415	137	no	0	no	12.2	
3	OH	area_code_408	84	no	0	yes	6.6	
4	OK	area_code_415	75	no	0	yes	10.1	
...	...	...	...	...	...	...	...	
4995	HI	area_code_408	50	yes	40	no	9.9	
4996	WV	area_code_415	152	no	0	no	14.7	
4997	DC	area_code_415	61	no	0	no	13.6	
4998	DC	area_code_510	109	no	0	no	8.5	
4999	VT	area_code_415	86	yes	34	no	9.3	

5000 rows × 23 columns



## Dropping Already merged columns

In [55]:

```
df.drop(['day.calls'],axis=1,inplace=True)
```

In [56]:

```
df.drop(['eve.calls'],axis=1,inplace=True)
```

In [57]:

```
df.drop(['night.calls'],axis=1,inplace=True)
```

In [58]:

```
df.drop(['day.charge'],axis=1,inplace=True)
```

In [59]:

```
df.drop(['eve.charge'],axis=1,inplace=True)
```

In [60]:

```
df.drop(['night.charge'],axis=1,inplace=True)
```

In [61]:

```
df.drop(['day.mins'],axis=1,inplace=True)
```

In [62]:

```
df.drop(['eve.mins'],axis=1,inplace=True)
```

In [63]:

```
df.drop(['night.mins'],axis=1,inplace=True)
```

## Encoding categorical columns

In [64]:

```
from sklearn.preprocessing import LabelEncoder
```

In [65]:

```
le=LabelEncoder()
```

In [66]:

```
for i in cat_feature:
    df[i]=le.fit_transform(df[i])
```

# Splitting data into x and y

In [67]:

```
x=df.drop('churn',axis=1)
```

In [68]:

x

Out[68]:

	state	area.code	account.length	voice.plan	voice.messages	intl.plan	intl.mins	intl.ca
0	16	1	128	1	25	0	10.0	
1	35	1	107	1	26	0	13.7	
2	31	1	137	0	0	0	12.2	
3	35	0	84	0	0	1	6.6	
4	36	1	75	0	0	1	10.1	
...	...	...	...	...	...	...	...	...
4995	11	0	50	1	40	0	9.9	
4996	49	1	152	0	0	0	14.7	
4997	7	1	61	0	0	0	13.6	
4998	7	2	109	0	0	0	8.5	
4999	46	1	86	1	34	0	9.3	

5000 rows × 13 columns





In [69]:

```
y=df['churn']  
y
```

Out[69]:

```
0      0  
1      0  
2      0  
3      0  
4      0  
..  
4995   0  
4996   1  
4997   0  
4998   0  
4999   0  
Name: churn, Length: 5000, dtype: int32
```

## Scaling Data

In [70]:

```
from sklearn.preprocessing import StandardScaler
```

In [71]:

```
sc=StandardScaler()
```

In [72]:

```
x=sc.fit_transform(x)
```

## Splitting data into train and test

In [73]:

```
from sklearn.model_selection import train_test_split
```

In [74]:

```
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.20,random_state=1)
```

**As we have imabalanced data we will do oversampling by smote as it will balanced out our data and we'll won't face data loss and there will be no more duplicate values in our data.**

In [75]:

```
from imblearn.over_sampling import SMOTE  
os = SMOTE(sampling_strategy=1.0)
```

In [76]:

```
xtrain_os, ytrain_os=os.fit_resample(xtrain,ytrain)
```

In [77]:

```
xtrain_os.shape
```

Out[77]:

```
(6862, 13)
```

In [78]:

```
xtest.shape
```

Out[78]:

```
(1000, 13)
```

In [79]:

```
ytrain_os.shape
```

Out[79]:

```
(6862,)
```

In [80]:

```
ytest.shape
```

Out[80]:

```
(1000,)
```

## Logistic Regression

In [81]:

```
from sklearn.linear_model import LogisticRegression
```

In [82]:

```
lr=LogisticRegression()
```

In [83]:

```
lr.fit(xtrain_os,ytrain_os)
```

Out[83]:

```
▼ LogisticRegression  
LogisticRegression()
```

In [84]:

```
y_pred_train=lr.predict(xtrain_os)  
y_pred_test=lr.predict(xtest)
```

In [85]:

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

## Accuracy Score of Logistic Regression

In [86]:

```
print("Train Data")  
print(accuracy_score(ytrain_os,y_pred_train))  
print("Test Data")  
print(accuracy_score(ytest,y_pred_test))
```

```
Train Data  
0.7892742640629554  
Test Data  
0.791
```

## Classification Report of Logistic Regression

In [87]:

```
print("Train Data")
print(classification_report(ytrain_os,y_pred_train))
print("Test Data")
print(classification_report(ytest,y_pred_test))
```

Train Data

	precision	recall	f1-score	support
0	0.80	0.77	0.79	3431
1	0.78	0.81	0.79	3431
accuracy			0.79	6862
macro avg	0.79	0.79	0.79	6862
weighted avg	0.79	0.79	0.79	6862

Test Data

	precision	recall	f1-score	support
0	0.95	0.80	0.87	862
1	0.37	0.73	0.49	138
accuracy			0.79	1000
macro avg	0.66	0.77	0.68	1000
weighted avg	0.87	0.79	0.82	1000

In [88]:

```
print("Train Data")
print(confusion_matrix(ytrain_os,y_pred_train))
print("Test Data")
print(confusion_matrix(ytest,y_pred_test))
```

Train Data

```
[[2649 782]
 [ 664 2767]]
```

Test Data

```
[[690 172]
 [ 37 101]]
```

## Decision Tree

In [89]:

```
from sklearn.tree import DecisionTreeClassifier
```

In [90]:

```
dt=DecisionTreeClassifier()
```

In [91]:

```
dt.fit(xtrain_os,ytrain_os)
```

Out[91]:

```
▼ DecisionTreeClassifier
DecisionTreeClassifier()
```

In [92]:

```
def my_model(clf):
    clf.fit(xtrain_os,ytrain_os)
    y_train_pred=clf.predict(xtrain_os)
    y_test_pred=clf.predict(xtest)
    print('Train Data')
    print(classification_report(ytrain_os,y_train_pred))
    print('Test Data')
    print(classification_report(ytest,y_test_pred))
```

In [93]:

```
my_model(dt)
```

Train Data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3431
1	1.00	1.00	1.00	3431
accuracy			1.00	6862
macro avg	1.00	1.00	1.00	6862
weighted avg	1.00	1.00	1.00	6862

Test Data

	precision	recall	f1-score	support
0	0.97	0.94	0.96	862
1	0.69	0.82	0.75	138
accuracy			0.93	1000
macro avg	0.83	0.88	0.85	1000
weighted avg	0.93	0.93	0.93	1000

## Hyper parameter tuning for Decision Tree

In [94]:

```
param_grid={'criterion': ['gini', 'entropy'],
            'splitter': ['best', 'random'],
            'max_depth': list(range(3, 51)),
            'min_samples_split': list(range(2, 50)),
            'min_samples_leaf': list(range(1, 50)),
            'max_features': ['auto', 'log2', None]
}
```

In [95]:

```
from sklearn.model_selection import RandomizedSearchCV
```

In [ ]:


In [96]:

```
clf=RandomizedSearchCV(dt,param_distributions=param_grid,n_iter=10,scoring='f1',n_jobs=-
```

In [97]:

```
clf.fit(xtrain_os,ytrain_os)
```

Out[97]:



In [98]:

```
clf.best_params_
```

Out[98]:

```
{'splitter': 'random',
 'min_samples_split': 3,
 'min_samples_leaf': 1,
 'max_features': None,
 'max_depth': 41,
 'criterion': 'gini'}
```

In [99]:

```
dt1=DecisionTreeClassifier(splitter='best',min_samples_split=46,min_samples_leaf=17,cri
```

In [100]:

```
my_model(dt1)
```

Train Data

	precision	recall	f1-score	support
0	0.94	0.97	0.95	3431
1	0.97	0.94	0.95	3431
accuracy			0.95	6862
macro avg	0.95	0.95	0.95	6862
weighted avg	0.95	0.95	0.95	6862

Test Data

	precision	recall	f1-score	support
0	0.97	0.96	0.96	862
1	0.75	0.81	0.78	138
accuracy			0.94	1000
macro avg	0.86	0.88	0.87	1000
weighted avg	0.94	0.94	0.94	1000

## Random Forest Classifier

In [101]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [102]:

```
rf=RandomForestClassifier()
```

In [103]:

```
my_model(rf)
```

Train Data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3431
1	1.00	1.00	1.00	3431
accuracy			1.00	6862
macro avg	1.00	1.00	1.00	6862
weighted avg	1.00	1.00	1.00	6862

Test Data

	precision	recall	f1-score	support
0	0.97	1.00	0.98	862
1	0.97	0.80	0.88	138
accuracy			0.97	1000
macro avg	0.97	0.90	0.93	1000
weighted avg	0.97	0.97	0.97	1000

## Hyper Parameter Tunning for Random Forest Classifier

In [104]:

```
param_grid1={'criterion': ['gini', 'entropy'],
             'max_depth': list(range(3, 51)),
             'min_samples_split': list(range(2, 50)),
             'min_samples_leaf': list(range(1, 50)),
             'max_features': ['auto', 'log2', None]}
}
```


In [105]:

```
clf1=RandomizedSearchCV(rf,param_distributions=param_grid1,n_iter=10,scoring='f1',n_jobs
```

In [106]:

```
clf1.fit(xtrain_os,ytrain_os)
```

Out[106]:





In [107]:

```
clf1.best_params_
```

Out[107]:

```
{'min_samples_split': 17,
 'min_samples_leaf': 6,
 'max_features': 'log2',
 'max_depth': 13,
 'criterion': 'entropy'}
```

In [108]:

```
rf1=RandomForestClassifier(min_samples_split=44,min_samples_leaf=1,criterion='entropy',m
```

In [109]:

```
my_model(rf1)
```

Train Data

	precision	recall	f1-score	support
0	0.95	0.99	0.97	3431
1	0.99	0.95	0.97	3431
accuracy			0.97	6862
macro avg	0.97	0.97	0.97	6862
weighted avg	0.97	0.97	0.97	6862

Test Data

	precision	recall	f1-score	support
0	0.97	0.99	0.98	862
1	0.91	0.83	0.87	138
accuracy			0.96	1000
macro avg	0.94	0.91	0.92	1000
weighted avg	0.96	0.96	0.96	1000

## AdaBoost Classifier

In [110]:

```
from sklearn.ensemble import AdaBoostClassifier
```

In [111]:

```
ada=AdaBoostClassifier()
```

In [112]:

my\_model(ada)

Train Data

	precision	recall	f1-score	support
0	0.90	0.93	0.91	3431
1	0.92	0.90	0.91	3431
accuracy			0.91	6862
macro avg	0.91	0.91	0.91	6862
weighted avg	0.91	0.91	0.91	6862

Test Data

	precision	recall	f1-score	support
0	0.95	0.93	0.94	862
1	0.60	0.69	0.64	138
accuracy			0.89	1000
macro avg	0.78	0.81	0.79	1000
weighted avg	0.90	0.89	0.90	1000

## Hyper parameter tuning for AdaBoost classifier


In [113]:

```
param_grid2={'n_estimators':[20,50,70,100,120],
             'learning_rate':(0.1,0.01,0.001,1),
             }
```

In [114]:

```
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error
ran_clf=RandomizedSearchCV(ada,param_distributions=param_grid2,cv=10,n_jobs=-1)
ran_clf.fit(xtrain_os,ytrain_os)
```

Out[114]:



In [115]:

ran\_clf.best\_params\_

Out[115]:

{'n\_estimators': 120, 'learning\_rate': 1}

In [116]:

```
ada1=AdaBoostClassifier(n_estimators= 50,learning_rate= 1)
```

In [117]:

```
my_model(ada1)
```

Train Data

	precision	recall	f1-score	support
0	0.90	0.93	0.91	3431
1	0.92	0.90	0.91	3431
accuracy			0.91	6862
macro avg	0.91	0.91	0.91	6862
weighted avg	0.91	0.91	0.91	6862

Test Data

	precision	recall	f1-score	support
0	0.95	0.93	0.94	862
1	0.60	0.69	0.64	138
accuracy			0.89	1000
macro avg	0.78	0.81	0.79	1000
weighted avg	0.90	0.89	0.90	1000

## Support Vector Machine Classifier

In [118]:

```
from sklearn.svm import SVC
```

In [119]:

```
svc=SVC()
```

In [120]:

```
my_model(svc)
```

Train Data

	precision	recall	f1-score	support
0	0.91	0.95	0.93	3431
1	0.94	0.90	0.92	3431
accuracy			0.93	6862
macro avg	0.93	0.93	0.93	6862
weighted avg	0.93	0.93	0.93	6862

Test Data

	precision	recall	f1-score	support
0	0.96	0.95	0.96	862
1	0.71	0.78	0.74	138
accuracy			0.93	1000
macro avg	0.84	0.86	0.85	1000
weighted avg	0.93	0.93	0.93	1000

## Gradient Boosting Classifier

In [121]:

```
from sklearn.ensemble import GradientBoostingClassifier
```

In [122]:

```
gb=GradientBoostingClassifier()
```

In [123]:

my\_model(gb)

Train Data

	precision	recall	f1-score	support
0	0.94	0.99	0.97	3431
1	0.99	0.94	0.97	3431
accuracy			0.97	6862
macro avg	0.97	0.97	0.97	6862
weighted avg	0.97	0.97	0.97	6862

Test Data

	precision	recall	f1-score	support
0	0.97	0.99	0.98	862
1	0.95	0.82	0.88	138
accuracy			0.97	1000
macro avg	0.96	0.91	0.93	1000
weighted avg	0.97	0.97	0.97	1000

## Hyper Parameter tuning for Gradient Boosting


In [124]:

```
gb_grid={'n_estimators':[50,100,120,150],
        'learning_rate':(0.1,0.01,0.001)
        }
```

In [125]:

```
GB_clf1=RandomizedSearchCV(gb,param_distributions=gb_grid,cv=10,n_jobs=-1)
GB_clf1.fit(xtrain_os,ytrain_os)
```

Out[125]:



In [126]:

GB\_clf1.best\_params\_

Out[126]:

{'n\_estimators': 120, 'learning\_rate': 0.1}

In [127]:

```
gb1=GradientBoostingClassifier(n_estimators= 150,learning_rate=0.1)
```

In [128]:

```
my_model(gb1)
```

Train Data

	precision	recall	f1-score	support
0	0.96	1.00	0.98	3431
1	1.00	0.96	0.98	3431
accuracy			0.98	6862
macro avg	0.98	0.98	0.98	6862
weighted avg	0.98	0.98	0.98	6862

Test Data

	precision	recall	f1-score	support
0	0.97	0.99	0.98	862
1	0.96	0.81	0.88	138
accuracy			0.97	1000
macro avg	0.96	0.90	0.93	1000
weighted avg	0.97	0.97	0.97	1000

## Cross Validation

### 1) K-Fold Cross Validation

In [129]:

```
from sklearn.model_selection import KFold
model=GradientBoostingClassifier(random_state=42)
kfold_validation=KFold(10)
import numpy as np
from sklearn.model_selection import cross_val_score
results=cross_val_score(model,x,y,cv=kfold_validation)
print(results)
print(round(np.mean(results),2))
```

```
[0.97 0.97 0.968 0.97 0.972 0.956 0.986 0.96 0.97 0.986]
0.97
```

### 2) Stratified K-Fold Cross Validation

In [130]:

```
from sklearn.model_selection import StratifiedKFold
skfold=StratifiedKFold(n_splits=5)
model=GradientBoostingClassifier(random_state=42)
scores=cross_val_score(model,x,y,cv=skfold)
print(round(np.mean(scores),2))
```

0.97

## Model Deployment

In [131]:

```
final_df=df[['voice.plan','voice.messages','intl.plan','intl.mins','intl.calls','intl.ch
final_df
```

Out[131]:

	voice.plan	voice.messages	intl.plan	intl.mins	intl.calls	intl.charge	customer.calls	Tc
0	1	25	0	10.0	3	2.70	1	
1	1	26	0	13.7	3	3.70	1	
2	0	0	0	12.2	5	3.29	0	
3	0	0	1	6.6	7	1.78	2	
4	0	0	1	10.1	3	2.73	3	
...	...	...	...	...	...	...	...	
4995	1	40	0	9.9	5	2.67	2	
4996	0	0	0	14.7	2	3.97	3	
4997	0	0	0	13.6	4	3.67	1	
4998	0	0	0	8.5	6	2.30	0	
4999	1	34	0	9.3	16	2.51	0	

5000 rows × 11 columns

In [132]:

```
final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   voice.plan            5000 non-null   int32   
 1   voice.messages        5000 non-null   int64   
 2   intl.plan             5000 non-null   int32   
 3   intl.mins             5000 non-null   float64  
 4   intl.calls            5000 non-null   int64   
 5   intl.charge           5000 non-null   float64  
 6   customer.calls        5000 non-null   int64   
 7   TotalDaysMins         5000 non-null   float64  
 8   TotalCharge           5000 non-null   float64  
 9   TotalDaysCalls        5000 non-null   int64   
10   churn                 5000 non-null   int32   
dtypes: float64(4), int32(3), int64(4)
memory usage: 371.2 KB
```

In [133]:

```
final_df.isnull().sum()
```

Out[133]:

```
voice.plan      0
voice.messages  0
intl.plan       0
intl.mins       0
intl.calls      0
intl.charge     0
customer.calls  0
TotalDaysMins   0
TotalCharge     0
TotalDaysCalls  0
churn           0
dtype: int64
```

In [134]:

```
x=final_df.drop(['churn'],axis=1)
y=final_df['churn']
```

In [149]:

```
from sklearn.model_selection import train_test_split
```

In [150]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_state=123)
```



In [151]:

```
from imblearn.over_sampling import SMOTE
os = SMOTE(sampling_strategy=1.0)
```

In [152]:

```
x_train_os, y_train_os=os.fit_resample(x_train,y_train)
```

In [153]:

```
x_train_os.shape,y_train_os.shape
```

Out[153]:

```
((6894, 10), (6894,))
```

In [154]:

```
x_test.shape,y_test.shape
```

Out[154]:

```
((1000, 10), (1000,))
```


In [155]:

```
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import Pipeline
from sklearn import metrics
```

In [159]:

```
pipe=Pipeline(steps=[('step1',LogisticRegression(solver='liblinear'))])
pipe.fit(x_train_os,y_train_os)
```

Out[159]:



In [160]:

```
y_pred=pipe.predict(x_test)
```

In [161]:

```
pipe.predict(x_test)[10]
```


Out[161]:

```
1
```

In [162]:

```
pipe1=Pipeline(steps=[('step1',GradientBoostingClassifier())])
pipe1.fit(x_train_os,y_train_os)
```

Out[162]:



In [163]:

```
pipe1.predict(x_test)[10]
```

Out[163]:

0

In [164]:

```
import pickle
pickle.dump(pipe,open('pipe.pkl','wb'))
```

In [165]:

```
final_df.sample(10)
```

Out[165]:

	voice.plan	voice.messages	intl.plan	intl.mins	intl.calls	intl.charge	customer.calls	T
4139	1	31	0	12.5	3	3.38	2	
873	0	0	0	9.3	1	2.51	1	
1048	0	0	0	7.7	3	2.08	0	
3309	0	0	1	12.0	4	3.24	4	
695	0	0	0	11.6	9	3.13	1	
4536	0	0	0	14.6	4	3.94	2	
2631	1	22	0	12.4	2	3.35	2	
3394	1	22	0	11.4	7	3.08	2	
2722	0	0	0	10.1	5	2.73	2	
2155	0	0	0	12.0	5	3.24	2	

In [ ]:

